


EdgeForge

Turns a trained **AI model** into **ready-to-run code for small, low-cost chips**, and checks that the chip version gives the same answers as the original.

THE BIG PICTURE What goes in, what comes out



No hand-written embedded code. One tool covers the whole path from data to a device-ready build.


WHY IT MATTERS The value in one line each

- **Cheaper hardware:** AI runs on chips costing a few dollars, no cloud or PC needed.
- **Faster delivery:** minutes from model to buildable code instead of weeks of manual porting.
- **Built-in trust:** every build is checked against the original model before it is handed over.
- **Usable by non-experts:** a point-and-click web page alongside the engineer's command line.

AT A GLANCE Today's reach


3 AI model formats accepted	5 chip profiles, 8-bit to 32-bit	256 B RAM on the smallest supported chip	3 output families: MCU, Arduino, FPGA	117 automated tests guarding quality
---------------------------------------	--	--	---	--

HOW IT WORKS One pipeline, two safety gates




The golden check runs the same sample inputs through the original model and the generated chip code, and compares the answers.

DESIGN CHOICE New chip = one settings file, no new code



Adding hardware is a configuration task. This keeps the cost of supporting new customer chips low.

STATUS Where the program stands



Bars show relative progress, not a schedule. Dashed portion is remaining scope.

NEXT Remaining scope

- Hosted multi-user service: accounts, secure hosting.
- One-click flashing to Arduino boards.
- FPGA: real board pin-mapping and larger models.

LIMITS What to know today

- Small and mid-size models only; no large language or speech models.
- Runs on one user's machine; not yet a shared service.
- Device code is a demo wiring; real sensor drivers are added per product.